Effects of Sparsity and the Activation Function on Sparse Autoencoders

Justin Johnson Bharath Ramsundar

Sparse Autoencoders

Autoencoders are neural networks that try to learn the identity function. Sparsity constraints force the learned representation to be nontrivial.

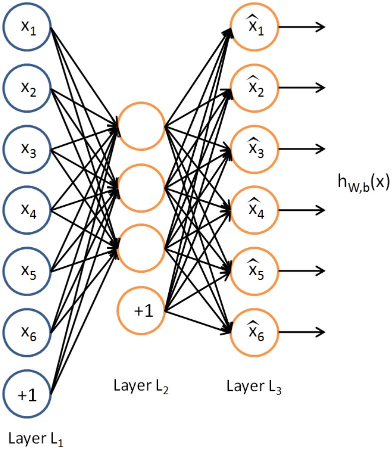

Figure: Autoencoder

Cost Function

Let the hidden layer have p units, and let $f : \mathbb{R} \to \mathbb{R}$ be a differentiable activation function. The neural network is parameterized by terms weights $W^{(1)} \in \mathbb{R}^{p \times n}$ and $W^{(2)} \in \mathbb{R}^{n \times p}$ and bias terms $b^{(1)} \in \mathbb{R}^p$ and $b^{(2)} \in \mathbb{R}^n$. The prediction on an input $x \in \mathbb{R}^n$ is

$$h_{W,b} = f(W^{(2)} f(W^{(1)} x + b^{(1)}) + b^{(2)})$$

Given training examples $x^{(1)}, \dots, x^{(m)} \in \mathbb{R}^n$ the objective function is

$$J(W, b) = \frac{1}{m} \sum_{i=1}^{m} \ell(h_{W,b}(x^{(i)}), x^{(i)}) + \lambda \psi(W, b) + \beta \sum_{j=1}^{p} \phi(\hat{\rho}_j)$$

where

$$\hat{\rho}_j = \frac{1}{m} \sum_{j=1}^{m} f\left((W_j^{(1)})^T x^{(i)} + b_j^{(1)} \right)$$

is the average activation of the jth hidden unit and $\ell : \mathbb{R}^n \times \mathbb{R}^n \to \mathbb{R}$ is a loss function. The function ψ is a regularizer, and ϕ is the sparsity function.

Pictures

Activation	L^1	L^2	KL
Sigmoid	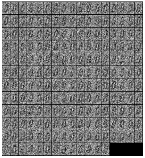	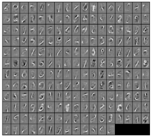	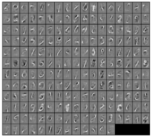
Identity